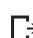


▼ Data Munging, Manipulation, Exploratory analysis using Pandas

```
import pandas as pd
import numpy as np
#Coding for importing csv files in Google colab
from google.colab import files
import io
uploaded = files.upload()
df = pd.read_csv(io.BytesIO(uploaded['loan.csv']))
# Read csv loan.csv into a pandas dataframe
# Take a look at the first few rows
print(df)
```

 No file chosen Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

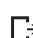
Saving loan.csv to loan.csv

	Loan_ID	Gender	Married	...	Loan_Amount_Term	Credit_History	Property_Area
0	LP001015	Male	Yes	...	360.0	1.0	Urban
1	LP001022	Male	Yes	...	360.0	1.0	Urban
2	LP001031	Male	Yes	...	360.0	1.0	Urban
3	LP001035	Male	Yes	...	360.0	NaN	Urban
4	LP001051	Male	No	...	360.0	1.0	Urban
..
362	LP002971	Male	Yes	...	360.0	1.0	Urban
363	LP002975	Male	Yes	...	360.0	1.0	Urban
364	LP002980	Male	No	...	360.0	NaN	Semiurban
365	LP002986	Male	Yes	...	360.0	1.0	Rural
366	LP002989	Male	No	...	180.0	1.0	Rural

[367 rows x 12 columns]

▼ To view the first 10 rows in the dataset

```
df.head(10)
df.columns
```

 Index(['Loan_ID', 'Gender', 'Married', 'Dependents', 'Education', 'Self_Employed', 'ApplicantIncome', 'CoapplicantIncome', 'LoanAmount', 'Loan_Amount_Term', 'Credit_History', 'Property_Area'], dtype='object')

▼ To calculate the statistical calculations for all numerical fields

```
df.describe()
```




	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	Credit_History
count	367.000000	367.000000	362.000000	361.000000	338.000000
mean	4805.599455	1569.577657	136.132597	342.537396	0.825444
std	4910.685399	2334.232099	61.366652	65.156643	0.380150
min	0.000000	0.000000	28.000000	6.000000	0.000000
25%	2864.000000	0.000000	100.250000	360.000000	1.000000
50%	3786.000000	1025.000000	125.000000	360.000000	1.000000

▼ Distribution analysis using EDA

Analysis on Application income alone using histogram

```
df['ApplicantIncome'].hist(bins=50)
```

☞ <matplotlib.axes._subplots.AxesSubplot at 0x7f1c6ccd2ba8>



▼ Analysis on Application income alone using boxplot

```
df.boxplot(column='ApplicantIncome')
```


☞

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f1c6f2f9978>
```

▼ Analysis on Application income and Education using boxplot

```
df.boxplot(column='ApplicantIncome', by = 'Education')
```


```
☐ <matplotlib.axes._subplots.AxesSubplot at 0x7f1c6cbf3eb8>
```



▼ Analysis on Application income and gender using boxplot

```
df.boxplot(column='ApplicantIncome', by = 'Gender')
```


```
☐ <matplotlib.axes._subplots.AxesSubplot at 0x7f1c6cd504e0>
```



▼ Analysis on Loan Amount alone using histogram

```
df['LoanAmount'].hist(bins=50)
```


↗ <matplotlib.axes._subplots.AxesSubplot at 0x7f1c6cb04208>



▼ Analysis on Gender alone using histogram

```
df['Gender'].hist(bins=50)
```


↗ <matplotlib.axes._subplots.AxesSubplot at 0x7f1c6cad7240>



▼ Analysis on Loan Amount alone using boxplot

```
df.boxplot(column='LoanAmount')
```


↗ <matplotlib.axes._subplots.AxesSubplot at 0x7f1c6c8bd5c0>



▼ Analysis on Loan Amount and gender using boxplot

```
df.boxplot(column='LoanAmount', by = 'Gender')
```

```
↳ <matplotlib.axes._subplots.AxesSubplot at 0x7f1c6ca97080>
```



▼ Categorical variable analysis

```
print ('Frequency Table for Credit History:')
temp1=df['Credit_History'].value_counts(ascending=True)
print(temp1)
```

```
print ('Frequency Table for Education:')
temp2=df['Education'].value_counts(ascending=True)
print(temp2)
```

```
↳ Frequency Table for Credit History:
0.0      59
1.0     279
Name: Credit_History, dtype: int64
Frequency Table for Education:
Not Graduate      84
Graduate         283
Name: Education, dtype: int64
```


▼ Applicants by Credit_History Analysis

```
import matplotlib.pyplot as plt
fig = plt.figure(figsize=(8,4))
```

```
#applicants by credit history
ax1 = fig.add_subplot(121)
ax1.set_xlabel('Credit_History')
ax1.set_ylabel('Count of Applicants')
```

```
ax1.set_xlabel('Count of Applicants')
ax1.set_title("Applicants by Credit_History")
temp1.plot(kind='bar')
```

☞ <matplotlib.axes._subplots.AxesSubplot at 0x7f0c9cb7a518>



Applicants by Credit_History Analysis and Applicants by Education Analysis both hand in hand


```
import matplotlib.pyplot as plt
fig = plt.figure(figsize=(8,4))

#applicants by credit history
ax1 = fig.add_subplot(121)
ax1.set_xlabel('Credit_History')
ax1.set_ylabel('Count of Applicants')
ax1.set_title("Applicants by Credit_History")
temp1.plot(kind='bar')
print('')

#applicants by education
ax2 = fig.add_subplot(122)
ax2.set_xlabel('Education')
ax2.set_ylabel('Count of Applicants')
ax2.set_title("Applicants by Education")
temp2.plot(kind='bar')
```

☞

<matplotlib.axes._subplots.AxesSubplot at 0x7f0c9c9fd6a0>



▼ Check missing values in the dataset

```
0      1      2      3      4      5      6      7      8      9      10      11      12      13      14      15      16      17      18      19      20      21      22      23      24      25      26      27      28      29      30      31      32      33      34      35      36      37      38      39      40      41      42      43      44      45      46      47      48      49      50      51      52      53      54      55      56      57      58      59      60      61      62      63      64      65      66      67      68      69      70      71      72      73      74      75      76      77      78      79      80      81      82      83      84      85      86      87      88      89      90      91      92      93      94      95      96      97      98      99
```

```
df.apply(lambda x: sum(x.isnull()),axis=0)
```

```
Loan_ID      0
Gender       11
Married      0
Dependents   10
Education    0
Self_Employed 23
ApplicantIncome 0
CoapplicantIncome 0
LoanAmount    5
Loan_Amount_Term 6
Credit_History 29
Property_Area 0
dtype: int64
```

▼ replacing missing loan amount with mean of the loanamount

```
df['LoanAmount'].fillna(df['LoanAmount'].mean(), inplace=True)
```

▼ viewing the data set

```
df
```

```
df
```

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	Coappli
0	LP001015	Male	Yes	0	Graduate	No	5720	
1	LP001022	Male	Yes	1	Graduate	No	3076	
2	LP001031	Male	Yes	2	Graduate	No	5000	
3	LP001035	Male	Yes	2	Graduate	No	2340	
4	LP001051	Male	No	0	Not Graduate	No	3276	

▼ once again checking empty values

363	LP002975	Male	Yes	0	Graduate	No	4158
-----	----------	------	-----	---	----------	----	------

```
df.apply(lambda x: sum(x.isnull()),axis=0)
```

```
↳ Loan_ID          0
   Gender          11
   Married         0
   Dependents      10
   Education        0
   Self_Employed   23
   ApplicantIncome  0
   CoapplicantIncome 0
   LoanAmount      0
   Loan_Amount_Term 6
   Credit_History   29
   Property_Area    0
   dtype: int64
```

▼ checking Self_Employed

```
df['Self_Employed'].value_counts()
```

```
↳ No      330
   Yes     37
   Name: Self_Employed, dtype: int64
```

▼ As No is dominating, replacing the empty values with No

```
df['Self_Employed'].fillna('No',inplace=True)
```

▼ checking Self_Employed once again

```
df['Self_Employed'].value_counts()
```



```
↳ No      330
   Yes     37
   Name: Self_Employed, dtype: int64
```

▼ checking Dependents

```
df['Dependents'].value_counts()
```

```
↳ 0      210
   2       59
   1       58
   3+      40
   Name: Dependents, dtype: int64
```

▼ As 0 is dominating , replace empty values with 0

```
df['Dependents'].fillna('0',inplace=True)
```

▼ once again checking Dependents

```
df['Dependents'].value_counts()
```

```
↳ 0      210
   2       59
   1       58
   3+      40
   Name: Dependents, dtype: int64
```

▼ once again checking empty values

```
df.apply(lambda x: sum(x.isnull()),axis=0)
```

```
↳
```

```
Loan_ID      0
Gender       0
```

▼ checking Gender

```
ApplicantIncome      0
```

```
df['Gender'].value_counts()
```

```
↵ Male      297
   Female    70
   Name: Gender, dtype: int64
```

▼ male is dominated with 80% so replace empty values with Male

```
df['Gender'].fillna('Male',inplace=True)
```

▼ once again checking Gender

```
df['Gender'].value_counts()
```

```
↵ Male      297
   Female    70
   Name: Gender, dtype: int64
```

▼ once again checking empty values

```
df.apply(lambda x: sum(x.isnull()),axis=0)
```

```
↵ Loan_ID      0
   Gender      0
   Married     0
   Dependents   0
   Education    0
   Self_Employed 0
   ApplicantIncome 0
   CoapplicantIncome 0
   LoanAmount   0
   Loan_Amount_Term 0
   Credit_History 0
   Property_Area 0
   dtype: int64
```

▼ checking Loan_Amount_Term

```
df['Loan_Amount_Term'].value_counts()
```

```
↳ 360.0    317
   180.0     22
   480.0      8
   300.0      7
   240.0      4
    84.0      3
     6.0      1
   120.0      1
    36.0      1
   350.0      1
    12.0      1
    60.0      1
Name: Loan_Amount_Term, dtype: int64
```

As Loan_Amount_Term=360 is dominating, replace empty values with 360

```
df['Loan_Amount_Term'].fillna(360.0,inplace=True)
```

checking Loan_Amount_Term

```
df['Loan_Amount_Term'].value_counts()
```

```
↳ 360.0    317
   180.0     22
   480.0      8
   300.0      7
   240.0      4
    84.0      3
     6.0      1
   120.0      1
    36.0      1
   350.0      1
    12.0      1
    60.0      1
Name: Loan_Amount_Term, dtype: int64
```

once again checking empty values

```
df.apply(lambda x: sum(x.isnull()),axis=0)
```

```
↳
```

Loan_ID	0
Gender	0
Married	0
Dependents	0
Education	0
Self_Employed	0
ApplicantIncome	0
CoapplicantIncome	0
LoanAmount	0
Loan_Amount_Term	0
Credit_History	0

▼ checking Credit_History

```
df['Credit_History'].value_counts()
```

```

1.0    308
0.0     59
Name: Credit_History, dtype: int64

```

▼ yes (1.0) is dominating

```
df['Credit_History'].fillna(1.0,inplace=True)
```

▼ once again checking empty values

```
df.apply(lambda x: sum(x.isnull()),axis=0)
```

```

Loan_ID      0
Gender        0
Married       0
Dependents    0
Education     0
Self_Employed 0
ApplicantIncome 0
CoapplicantIncome 0
LoanAmount    0
Loan_Amount_Term 0
Credit_History 0
Property_Area 0
dtype: int64

```


▼ Finally all missing values are clear

Then go to the next phase of normalization

how to treat for extreme values in distribution of LoanAmount and ApplicantIncome

```
df['LoanAmount'].hist(bins=10)
```


↗ <matplotlib.axes._subplots.AxesSubplot at 0x7f1c6c725198>



creating LoanAmount_log column to treat outliers and extreme values

```
df['LoanAmount_log'] = np.log(df['LoanAmount'])  
df['LoanAmount_log'].hist(bins=20)
```

↗ <matplotlib.axes._subplots.AxesSubplot at 0x7f1c6c76b4e0>



The normalized data set with artificial field LoanAmount_log

df



	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	Coappli
0	LP001015	Male	Yes	0	Graduate	No	5720	
1	LP001022	Male	Yes	1	Graduate	No	3076	
2	LP001031	Male	Yes	2	Graduate	No	5000	
3	LP001035	Male	Yes	2	Graduate	No	2340	
4	LP001051	Male	No	0	Not Graduate	No	3276	
...
362	LP002971	Male	Yes	3+	Not Graduate	Yes	4009	
363	LP002975	Male	Yes	0	Graduate	No	4158	
364	LP002980	Male	No	0	Graduate	No	3250	
365	LP002986	Male	Yes	0	Graduate	No	5000	
366	LP002989	Male	No	0	Graduate	Yes	9200	

367 rows × 13 columns

so far we have treated an untreated data set from repository. But We have some treated data set with relevant key column which indicates the result of the model.

Loading already treated one such loan approval dataset from repository to predict the loan approval with Loan_status field

```
import pandas as pd
url = "https://raw.githubusercontent.com/callxpert/datasets/master/Loan-applicant-details.csv"
names = ['Loan_ID', 'Gender', 'Married', 'Dependents', 'Education', 'Self_Employed', 'ApplicantIncome', 'CoapplicantIncome', 'Loan_Status']
df = pd.read_csv(url, names=names)
print(df)
```



```

      Loan_ID  Gender Married  ... Credit_History Property_Area Loan_Status
0      LP001003   Male    Yes  ...              1         Rural           N

df.apply(lambda x: sum(x.isnull()),axis=0)

```

```

☞ Loan_ID      0
   Gender      0
   Married     0
   Dependents  0
   Education   0
   Self_Employed 0
   ApplicantIncome 0
   CoapplicantIncome 0
   LoanAmount     0
   Loan_Amount_Term 0
   Credit_History  0
   Property_Area   0
   Loan_Status     0
   dtype: int64

```

▼ Lets take a peek at the data

```
print(df.head(20))
```

```

☞      Loan_ID  Gender Married  ... Credit_History Property_Area Loan_Status
0      LP001003   Male    Yes  ...              1         Rural           N
1      LP001005   Male    Yes  ...              1         Urban           Y
2      LP001006   Male    Yes  ...              1         Urban           Y
3      LP001008   Male    No   ...              1         Urban           Y
4      LP001011   Male    Yes  ...              1         Urban           Y
5      LP001013   Male    Yes  ...              1         Urban           Y
6      LP001014   Male    Yes  ...              0      Semiurban           N
7      LP001018   Male    Yes  ...              1         Urban           Y
8      LP001020   Male    Yes  ...              1      Semiurban           N
9      LP001024   Male    Yes  ...              1         Urban           Y
10     LP001028   Male    Yes  ...              1         Urban           Y
11     LP001029   Male    No   ...              1         Rural           N
12     LP001030   Male    Yes  ...              1         Urban           Y
13     LP001032   Male    No   ...              1         Urban           Y
14     LP001036  Female    No   ...              0         Urban           N
15     LP001038   Male    Yes  ...              1         Rural           N
16     LP001043   Male    Yes  ...              0         Urban           N
17     LP001046   Male    Yes  ...              1         Urban           Y
18     LP001047   Male    Yes  ...              0      Semiurban           N
19     LP001066   Male    Yes  ...              1      Semiurban           Y

```

```
[20 rows x 13 columns]
```

▼ Lets load the required libraries for our analysis

```

#Load libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn import model_selection
from sklearn.metrics import accuracy_score

```

```

from sklearn.linear_model import LogisticRegression
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier, export_graphviz
from sklearn.model_selection import train_test_split

```

Steps involved in this machine learning project

Following are the steps involved in creating a well-defined ML project**

Understand and define the problem ,Analyse and prepare the data ,Apply the algorithms ,Reduce the errors ,Predict the result

```

from sklearn.preprocessing import LabelEncoder
var_mod = ['Gender', 'Married', 'Dependents', 'Education', 'Self_Employed', 'Property_Area', 'Loan_Status']
le = LabelEncoder()
for i in var_mod:
    df[i] = le.fit_transform(df[i])

```

sklearn requires all inputs to be numeric, we should convert all our categorical variables into numeric by encoding the categories. This can be done using the above code:

Splitting the Data set: As we have seen already, In Machine learning we have two kinds of datasets

Training dataset - used to train our model

Testing dataset - used to test if our model is making accurate predictions

Our dataset has 480 records. We are going to use 80% of it for training the model and 20% of the records to evaluate our model. copy paste the below commands to prepare our data sets

```

array = df.values
X = array[:,6:11]
Y = array[:,12]
Y=Y.astype('int')
x_train, x_test, y_train, y_test = model_selection.train_test_split(X, Y, test_size=0.2, random_stat
df.columns

```




```
Index(['Loan_ID', 'Gender', 'Married', 'Dependents', 'Education',  
      'Loan_Amount_Term', 'Credit_History', 'Property_Area',  
      'Self_Employed', 'ApplicantIncome', 'CoapplicantIncome',  
      'Loan_Amount', 'Loan_Term'],  
      dtype='object', name='index')
```

Evaluating the model and training the Model with

▼ 'ApplicantIncome', 'CoapplicantIncome',
'LoanAmount','Loan_Amount_Term', 'Credit_History'

ML model 1

Logistic Regression : Logistic Regression is a classification algorithm. It is used to predict a binary outcome (1 / 0, Yes / No, True / False) given a set of independent variables. To represent binary / categorical outcome, we use dummy variables

```
model = LogisticRegression()  
model.fit(x_train,y_train)  
predictions = model.predict(x_test)  
print(accuracy_score(y_test, predictions))
```

☞ 0.7708333333333334

Decision tree : Decision tree is a type of supervised learning algorithm (having a pre-defined target variable) that is mostly used in classification problems. It works for both categorical and continuous input and output variables

```
model = DecisionTreeClassifier()  
model.fit(x_train,y_train)  
predictions = model.predict(x_test)  
print(accuracy_score(y_test, predictions))
```

☞ 0.6354166666666666

Random forest : Random forests or random decision forests are an ensemble learning method for classification, regression and other tasks, that operate by constructing a multitude of decision trees at training time and outputting the class that is the mode of the classes (classification) or mean prediction (regression) of the individual trees

```
model = RandomForestClassifier(n_estimators=100)  
model.fit(x_train,y_train)  
predictions = model.predict(x_test)  
print(accuracy_score(y_test, predictions))
```

☞ 0.75

Evaluating the model and training the Model with 'Married',
▼ 'Dependents', 'Education','Self_Employed', 'ApplicantIncome'

ML model 2

```
array = df.values
X = array[:,2:6]
Y = array[:,12]
Y=Y.astype('int')
x_train, x_test, y_train, y_test = model_selection.train_test_split(X, Y, test_size=0.2, random_state=42)
df.columns
```

```
↳ Index(['Loan_ID', 'Gender', 'Married', 'Dependents', 'Education',
        'Self_Employed', 'ApplicantIncome', 'CoapplicantIncome', 'LoanAmount',
        'Loan_Amount_Term', 'Credit_History', 'Property_Area', 'Loan_Status'],
       dtype='object')
```

```
model = LogisticRegression()
model.fit(x_train,y_train)
predictions = model.predict(x_test)
print(accuracy_score(y_test, predictions))
```

```
↳ 0.6354166666666666
```

```
model = DecisionTreeClassifier()
model.fit(x_train,y_train)
predictions = model.predict(x_test)
print(accuracy_score(y_test, predictions))
```

```
↳ 0.6145833333333334
```

```
model = RandomForestClassifier(n_estimators=100)
model.fit(x_train,y_train)
predictions = model.predict(x_test)
print(accuracy_score(y_test, predictions))
```

```
↳ 0.625
```

Double-click (or enter) to edit

```
array = df.values
X = array[:,5:12]
Y = array[:,12]
Y=Y.astype('int')
x_train, x_test, y_train, y_test = model_selection.train_test_split(X, Y, test_size=0.2, random_state=42)
df.columns
```

```
↳ Index(['Loan_ID', 'Gender', 'Married', 'Dependents', 'Education',
        'Self_Employed', 'ApplicantIncome', 'CoapplicantIncome', 'LoanAmount',
        'Loan_Amount_Term', 'Credit_History', 'Property_Area', 'Loan_Status'],
       dtype='object')
```

Evaluating the model and training the Model with 'Self_Employed',
'ApplicantIncome', 'CoapplicantIncome',
'LoanAmount','Loan_Amount_Term', 'Credit_History',
'Property_Area'

```
array = df.values
X = array[:,5:12]
Y = array[:,12]
Y=Y.astype('int')
x_train, x_test, y_train, y_test = model_selection.train_test_split(X, Y, test_size=0.2, random_state=42)
df.columns
```

```
Index(['Loan_ID', 'Gender', 'Married', 'Dependents', 'Education',
       'Self_Employed', 'ApplicantIncome', 'CoapplicantIncome', 'LoanAmount',
       'Loan_Amount_Term', 'Credit_History', 'Property_Area', 'Loan_Status'],
      dtype='object')
```

Double-click (or enter) to edit

```
model = LogisticRegression()
model.fit(x_train,y_train)
predictions = model.predict(x_test)
print(accuracy_score(y_test, predictions))
```

```
0.7708333333333334
```

Double-click (or enter) to edit

```
model = DecisionTreeClassifier()
model.fit(x_train,y_train)
predictions = model.predict(x_test)
print(accuracy_score(y_test, predictions))
```

```
0.65625
```

Double-click (or enter) to edit

```
model = RandomForestClassifier(n_estimators=100)
model.fit(x_train,y_train)
predictions = model.predict(x_test)
print(accuracy_score(y_test, predictions))
```

```
0.7395833333333334
```

