

## ▼ Data Munging, Manipulation, Exploratory analysis using Pandas

```
import pandas as pd
import numpy as np
#Coding for importing csv files in Google colab
from google.colab import files
import io
uploaded = files.upload()
df = pd.read_csv(io.BytesIO(uploaded['loan.csv']))
# Read csv loan.csv into a pandas dataframe
# Take a look at the first few rows
print(df)
```



Choose Files

No file chosen

Upload widget is only available when the cell has been executed in the current browser session. Please rerun this cell to enable.

Saving loan.csv to loan.csv

	Loan_ID	Gender	Married	...	Loan_Amount_Term	Credit_History	Property_Area
0	LP001015	Male	Yes	...	360.0	1.0	Urban
1	LP001022	Male	Yes	...	360.0	1.0	Urban
2	LP001031	Male	Yes	...	360.0	1.0	Urban
3	LP001035	Male	Yes	...	360.0	NaN	Urban
4	LP001051	Male	No	...	360.0	1.0	Urban
..	...	...	...	...	...	...	...
362	LP002971	Male	Yes	...	360.0	1.0	Urban
363	LP002975	Male	Yes	...	360.0	1.0	Urban
364	LP002980	Male	No	...	360.0	NaN	Semiurban
365	LP002986	Male	Yes	...	360.0	1.0	Rural
366	LP002989	Male	No	...	180.0	1.0	Rural

[367 rows x 12 columns]

## ▼ To view the first 10 rows in the dataset

```
df.head(10)
```

```
df.columns
```

```
Index(['Loan_ID', 'Gender', 'Married', 'Dependents', 'Education',  
      'Self_Employed', 'ApplicantIncome', 'CoapplicantIncome', 'LoanAmount',  
      'Loan_Amount_Term', 'Credit_History', 'Property_Area'],  
      dtype='object')
```

## ▼ To calculate the statistical calculations for all numerical fields

```
df.describe()
```

```

ApplicantIncome  CoapplicantIncome  LoanAmount  Loan_Amount_Term  Credit_History
count           367.000000          367.000000   362.000000          361.000000          338.000000
mean           4805.599455          1569.577657   136.132597          342.537396           0.825444
std           4910.685399          2334.232099    61.366652           65.156643           0.380150
min              0.000000           0.000000    28.000000           6.000000           0.000000
25%            2864.000000           0.000000   100.250000          360.000000          1.000000
50%            3786.000000          1025.000000   125.000000          360.000000          1.000000
75%            5060.000000          2430.500000   158.000000          360.000000          1.000000
max           72529.000000         24000.000000   550.000000          480.000000          1.000000
```


## ▼ Distribution analysis using EDA

Analysis on Application income alone using histogram

```
df['ApplicantIncome'].hist(bins=50)
```

```
df['ApplicantIncome'].hist(bins=50)
```

↳ <matplotlib.axes.\_subplots.AxesSubplot at 0x7f1c6ccd2ba8>



## ▼ Analysis on Application income alone using boxplot


```
df.boxplot(column='ApplicantIncome')
```

↳

## Analysis on Application income and Education using boxplot

```
df.boxplot(column='ApplicantIncome', by = 'Education')
```

↳ <matplotlib.axes.\_subplots.AxesSubplot at 0x7f1c6cbf3eb8>




## Analysis on Application income and gender using boxplot

```
df.boxplot(column='ApplicantIncome', by = 'Gender')
```

↳

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f1c6cd504e0>
```




## Analysis on Loan Amount alone using histogram



```
df['LoanAmount'].hist(bins=50)
```


```
<matplotlib.axes._subplots.AxesSubplot at 0x7f1c6cb04208>
```



## Analysis on Gender alone using histogram

```
df['Gender'].hist(bins=50)
```


↳ <matplotlib.axes.\_subplots.AxesSubplot at 0x7f1c6cad7240>



## ▼ Analysis on Loan Amount alone using boxplot

```
df.boxplot(column='LoanAmount')
```


↳ <matplotlib.axes.\_subplots.AxesSubplot at 0x7f1c6c8bd5c0>



## ▼ Analysis on Loan Amount and gender using boxplot

```
df.boxplot(column='LoanAmount', by = 'Gender')
```

```
↳ <matplotlib.axes._subplots.AxesSubplot at 0x7f1c6ca97080>
```



## ▼ Categorical variable analysis

```
print ('Frequency Table for Credit History:')  
temp1=df['Credit_History'].value_counts(ascending=True)  
print(temp1)
```

```
print ('Frequency Table for Education:')  
temp2=df['Education'].value_counts(ascending=True)  
print(temp2)
```

```
↳
```

Frequency Table for Credit History:

0.0      59

1.0      279

Name: Credit\_History, dtype: int64

Frequency Table for Education:

## ▼ Applicants by Credit\_History Analysis

```
import matplotlib.pyplot as plt
fig = plt.figure(figsize=(8,4))

#applicants by credit history
ax1 = fig.add_subplot(121)
ax1.set_xlabel('Credit_History')
ax1.set_ylabel('Count of Applicants')
ax1.set_title("Applicants by Credit_History")
temp1.plot(kind='bar')
```





<matplotlib.axes. subplots.AxesSubplot at 0x7f0c9cb7a518>

## Applicants by Credit\_History Analysis and Applicants by Education Analysis both hand in hand


```
import matplotlib.pyplot as plt
fig = plt.figure(figsize=(8,4))

#applicants by credit history
ax1 = fig.add_subplot(121)
ax1.set_xlabel('Credit_History')
ax1.set_ylabel('Count of Applicants')
ax1.set_title("Applicants by Credit_History")
temp1.plot(kind='bar')
print('')

#applicants by education
ax2 = fig.add_subplot(122)
ax2.set_xlabel('Education')
ax2.set_ylabel('Count of Applicants')
ax2.set_title("Applicants by Education")
temp2.plot(kind='bar')
```



<matplotlib.axes.\_subplots.AxesSubplot at 0x7f0c9c9fd6a0>



## ▼ Check missing values in the dataset

```
df.apply(lambda x: sum(x.isnull()),axis=0)
```

```
Loan_ID      0
Gender       11
Married      0
Dependents   10
Education    0
Self_Employed 23
ApplicantIncome 0
CoapplicantIncome 0
LoanAmount    5
Loan_Amount_Term 6
Credit_History 29
Property_Area 0
dtype: int64
```

## ▼ replacing missing loan amount with mean of the loanamount

```
df['LoanAmount'].fillna(df['LoanAmount'].mean(), inplace=True)
```

## ▼ viewing the data set

df

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	C
0	LP001015	Male	Yes	0	Graduate	No	5720	0	110.0	360.0	
1	LP001022	Male	Yes	1	Graduate	No	3076	1500	126.0	360.0	
2	LP001031	Male	Yes	2	Graduate	No	5000	1800	208.0	360.0	
3	LP001035	Male	Yes	2	Graduate	No	2340	2546	100.0	360.0	
4	LP001051	Male	No	0	Not Graduate	No	3276	0	78.0	360.0	
...	...	...	...	...	...	...	...	...	...	...	
362	LP002971	Male	Yes	3+	Not Graduate	Yes	4009	1777	113.0	360.0	
363	LP002975	Male	Yes	0	Graduate	No	4158	709	115.0	360.0	
364	LP002980	Male	No	0	Graduate	No	3250	1993	126.0	360.0	
365	LP002986	Male	Yes	0	Graduate	No	5000	2393	158.0	360.0	
366	LP002989	Male	No	0	Graduate	Yes	9200	0	98.0	180.0	

367 rows × 12 columns

once again checking empty values

```
df.apply(lambda x: sum(x.isnull()),axis=0)
```



Loan_ID	0
Gender	11
Married	0
Dependents	10
Education	0
Self_Employed	23
ApplicantIncome	0
CoapplicantIncome	0
LoanAmount	0

## ▼ checking Self\_Employed

```
df['Self_Employed'].value_counts()
```

```
↵ No      330  
   Yes      37  
   Name: Self_Employed, dtype: int64
```

## ▼ As No is dominating, replacing the empty values with No

```
df['Self_Employed'].fillna('No',inplace=True)
```

## ▼ checking Self\_Employed once again

```
df['Self_Employed'].value_counts()
```

```
↵ No      330  
   Yes      37  
   Name: Self_Employed, dtype: int64
```

## ▼ checking Dependents

```
df['Dependents'].value_counts()
```

```
☐➔  0      210  
    2       59  
    1       58  
    3+      40  
    Name: Dependents, dtype: int64
```

## ▼ As 0 is dominating , replace empty values with 0

```
df['Dependents'].fillna('0',inplace=True)
```

## ▼ once again checking Dependents

```
df['Dependents'].value_counts()
```

```
☐➔  0      210  
    2       59  
    1       58  
    3+      40  
    Name: Dependents, dtype: int64
```

## ▼ once again checking empty values

```
df.apply(lambda x: sum(x.isnull()),axis=0)
```

```
↳ Loan_ID      0
   Gender       0
   Married      0
   Dependents   0
   Education    0
   Self_Employed 0
   ApplicantIncome 0
   CoapplicantIncome 0
   LoanAmount   0
   Loan_Amount_Term 0
   Credit_History 0
   Property_Area 0
   dtype: int64
```

## ▼ checking Gender

```
df['Gender'].value_counts()
```

```
↳ Male      297
   Female    70
   Name: Gender, dtype: int64
```

## ▼ male is dominated with 80% so replace empty values with Male

```
df['Gender'].fillna('Male',inplace=True)
```

## ▼ once again checking Gender

```
df['Gender'].value_counts()
```

```
Male      297  
Female     70  
Name: Gender, dtype: int64
```

## ▼ once again checking empty values

```
df.apply(lambda x: sum(x.isnull()),axis=0)
```

```
Loan_ID      0  
Gender       0  
Married      0  
Dependents   0  
Education    0  
Self_Employed 0  
ApplicantIncome 0  
CoapplicantIncome 0  
LoanAmount   0  
Loan_Amount_Term 0  
Credit_History 0  
Property_Area 0  
dtype: int64
```

## ▼ checking Loan\_Amount\_Term

```
df['Loan_Amount_Term'].value_counts()
```

360.0	317
180.0	22
480.0	8
300.0	7
240.0	4
84.0	3
6.0	1
120.0	1
36.0	1
350.0	1
12.0	1
60.0	1

- As Loan\_Amount\_Term=360 is dominating, replace empty values with 360

```
df['Loan_Amount_Term'].fillna(360.0,inplace=True)
```

- checking Loan\_Amount\_Term

```
df['Loan_Amount_Term'].value_counts()
```

```

360.0    317
180.0     22
480.0      8
300.0      7
240.0      4
84.0       3
6.0        1
120.0      1
36.0       1
350.0      1
12.0       1
60.0       1
Name: Loan_Amount_Term, dtype: int64

```



## ▼ once again checking empty values

```
df.apply(lambda x: sum(x.isnull()),axis=0)
```

```
☞ Loan_ID          0
   Gender          0
   Married         0
   Dependents      0
   Education       0
   Self_Employed   0
   ApplicantIncome 0
   CoapplicantIncome 0
   LoanAmount      0
   Loan_Amount_Term 0
   Credit_History   0
   Property_Area    0
   dtype: int64
```

## ▼ checking Credit\_History

```
df['Credit_History'].value_counts()
```

```
☞ 1.0    308
   0.0     59
   Name: Credit_History, dtype: int64
```

## ▼ yes (1.0) is dominating

```
df['Credit_History'].fillna(1.0,inplace=True)
```

## ▼ once again checking empty values

```
df.apply(lambda x: sum(x.isnull()),axis=0)
```

```
↳ Loan_ID          0
   Gender           0
   Married          0
   Dependents       0
   Education        0
   Self_Employed    0
   ApplicantIncome  0
   CoapplicantIncome 0
   LoanAmount       0
   Loan_Amount_Term 0
   Credit_History   0
   Property_Area    0
   dtype: int64
```

## ▼ Finally all missing values are clear


Then go to the next phase of normalization

how to treat for extreme values in distribution of LoanAmount and ApplicantIncome

```
df['LoanAmount'].hist(bins=10)
```

```
↳
```


<matplotlib.axes.\_subplots.AxesSubplot at 0x7f1c6c725198>



- ▼ creating LoanAmount\_log column to treat outliers and extreme values

```
df['LoanAmount_log'] = np.log(df['LoanAmount'])  
df['LoanAmount_log'].hist(bins=20)
```

↳ <matplotlib.axes.\_subplots.AxesSubplot at 0x7f1c6c76b4e0>



- ▼ The normalized data set with artificial field LoanAmount\_log

↗

	Loan_ID	Gender	Married	Dependents	Education	Self_Employed	ApplicantIncome	CoapplicantIncome	LoanAmount	Loan_Amount_Term	C
0	LP001015	Male	Yes	0	Graduate	No	5720	0	110.0	360.0	
1	LP001022	Male	Yes	1	Graduate	No	3076	1500	126.0	360.0	
2	LP001031	Male	Yes	2	Graduate	No	5000	1800	208.0	360.0	
3	LP001035	Male	Yes	2	Graduate	No	2340	2546	100.0	360.0	
4	LP001051	Male	No	0	Not Graduate	No	3276	0	78.0	360.0	
...	...	...	...	...	...	...	...	...	...	...	
362	LP002971	Male	Yes	3+	Not Graduate	Yes	4009	1777	113.0	360.0	
363	LP002975	Male	Yes	0	Graduate	No	4158	709	115.0	360.0	
364	LP002980	Male	No	0	Graduate	No	3250	1993	126.0	360.0	
365	LP002986	Male	Yes	0	Graduate	No	5000	2393	158.0	360.0	
366	LP002989	Male	No	0	Graduate	Yes	9200	0	98.0	180.0	

367 rows × 13 columns

Loading already treated dataset from repository to predict the loan approval with Loan\_status field

```
url = "https://raw.githubusercontent.com/callxpert/datasets/master/Loan-applicant-details.csv"
names = ['Loan_ID','Gender','Married','Dependents','Education','Self_Employed','ApplicantIncome','CoapplicantIncome']
df = pd.read_csv(url, names=names)
print(df)
```

```

↳
   Loan_ID  Gender  Married  ...  Credit_History  Property_Area  Loan_Status
0    LP001003   Male    Yes    ...              1         Rural          N
1    LP001005   Male    Yes    ...              1         Urban          Y
2    LP001006   Male    Yes    ...              1         Urban          Y
3    LP001008   Male    No     ...              1         Urban          Y
4    LP001011   Male    Yes    ...              1         Urban          Y
..      ...      ...      ...      ...              ...          ...          ...
475  LP002978  Female    No     ...              1         Rural          Y
476  LP002979   Male    Yes    ...              1         Rural          Y
477  LP002983   Male    Yes    ...              1         Urban          Y
478  LP002984   Male    Yes    ...              1         Urban          Y
479  LP002990  Female    No     ...              0        Semiurban          N

```

[480 rows x 13 columns]

▼ let us check the null values in the new dataset downloaded from the repository

```
df.apply(lambda x: sum(x.isnull()),axis=0)
```

```

↳
Loan_ID          0
Gender           0
Married          0
Dependents       0
Education        0
Self_Employed    0
ApplicantIncome  0
CoapplicantIncome 0
LoanAmount       0
Loan_Amount_Term 0
Credit_History  0
Property_Area    0
Loan_Status      0
dtype: int64

```

▼ Lets take a peek at the data

```
print(dataset.head(20))
```

```
↳
```

	Loan_ID	Gender	Married	...	Credit_History	Property_Area	Loan_Status
0	LP001003	Male	Yes	...	1	Rural	N
1	LP001005	Male	Yes	...	1	Urban	Y
2	LP001006	Male	Yes	...	1	Urban	Y
3	LP001008	Male	No	...	1	Urban	Y
4	LP001011	Male	Yes	...	1	Urban	Y
5	LP001013	Male	Yes	...	1	Urban	Y
6	LP001014	Male	Yes	...	0	Semiurban	N
7	LP001018	Male	Yes	...	1	Urban	Y
8	LP001020	Male	Yes	...	1	Semiurban	N
9	LP001024	Male	Yes	...	1	Urban	Y
10	LP001028	Male	Yes	...	1	Urban	Y
11	LP001029	Male	No	...	1	Rural	N
12	LP001030	Male	Yes	...	1	Urban	Y
13	LP001032	Male	No	...	1	Urban	Y
14	LP001036	Female	No	...	0	Urban	N
15	LP001038	Male	Yes	...	1	Rural	N
16	LP001043	Male	Yes	...	0	Urban	N
17	LP001046	Male	Yes	...	1	Urban	Y
18	LP001047	Male	Yes	...	0	Semiurban	N
19	LP001066	Male	Yes	...	1	Semiurban	Y

```
[20 rows x 13 columns]
```

▼ Lets load the required libraries for analyzing our developed Model

```
#Load libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn import model_selection
from sklearn.metrics import accuracy_score
from sklearn.linear_model import LogisticRegression
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.neighbors import KNeighborsClassifier
```

```
from sklearn.neighbors import KNeighborsClassifier
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier, export_graphviz
from sklearn.model_selection import train_test_split
```

sklearn requires all inputs to be numeric, we should convert all our categorical variables into numeric by encoding the categories. This can be done using the following code:

```
from sklearn.preprocessing import LabelEncoder
var_mod = ['Gender', 'Married', 'Dependents', 'Education', 'Self_Employed', 'Property_Area', 'Loan_Status']
le = LabelEncoder()
for i in var_mod:
    dataset[i] = le.fit_transform(dataset[i])
```

**Splitting the Data set** As we have seen already, In Machine learning we have two kinds of datasets

**Training dataset** - used to train our model **Testing dataset** - used to test if our model is making accurate predictions Our dataset has 480 records. We are going to use 80% of it for training the model and 20% of the records to evaluate our model. copy paste the below commands to prepare our data sets

Though our dataset has lot of columns, we are only going to use the Income fields, loan amount, loan duration and credit history fields to train our model.

```
array = dataset.values
X = array[:,6:11]
Y = array[:,12]
Y=Y.astype('int')
x_train, x_test, y_train, y_test = model_selection.train_test_split(X, Y, test_size=0.2, random_state=7)
print(X)
print(Y)
```

```

↳ [[4583 1508.0 128 360 1]
    [3000 0.0 66 360 1]
    [2583 2358.0 120 360 1]
    ...
    [8072 240.0 253 360 1]
    [7583 0.0 187 360 1]
    [4583 0.0 133 360 0]]
[0 1 1 1 1 1 0 1 0 1 1 0 1 1 0 0 0 1 0 1 1 1 0 0 0 1 0 1 1 1 0 1 1 1 0 1
 1 1 0 0 0 1 1 0 1 1 1 1 0 0 0 0 1 1 0 1 1 1 1 0 0 0 0 1 0 1 1 1 1 1 1 1
 1 1 1 1 1 1 1 1 1 1 0 1 1 1 1 0 1 1 1 1 1 0 1 1 1 1 0 0 1 0 0 0 1 1 1 1
 1 1 0 1 0 1 0 0 1 1 1 1 1 0 0 1 1 0 1 0 1 0 1 0 1 1 0 1 0 0 1 0 1 1 0 1 1
 0 0 1 1 0 1 0 1 1 1 1 1 1 1 1 0 0 0 1 1 1 1 0 1 0 1 1 1 1 0 1 1 1 1 1 1 1
 1 0 1 1 1 0 1 1 1 1 0 0 1 1 0 1 0 0 0 1 0 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 0
 1 1 1 1 0 1 1 1 1 0 1 1 1 1 0 0 0 1 1 1 1 0 0 1 1 1 1 1 0 1 1 1 0 1 1 0 1
 1 1 1 1 0 0 1 1 1 1 1 1 0 1 0 1 1 0 0 1 0 1 1 1 0 0 1 0 1 1 1 0 1 1 0 1 1
 1 1 0 1 1 1 1 1 1 1 0 1 1 0 0 0 1 1 0 1 1 1 0 0 0 0 1 0 1 0 0 1 1 1 1 1 1
 0 1 1 1 1 0 1 1 1 0 1 1 1 1 1 1 0 0 0 1 1 1 1 0 1 0 1 1 0 1 0 0 1 0 1 1
 1 1 1 1 1 1 1 1 1 0 0 1 0 1 1 0 1 1 1 0 1 1 1 1 1 1 0 1 1 0 0 1 1 0 0 1
 1 1 1 1 1 1 1 1 1 0 1 1 1 0 1 1 0 1 1 0 1 1 0 0 1 1 1 1 1 1 1 1 0 1 0
 0 0 1 1 0 0 1 1 1 1 0 1 0 0 1 1 0 1 0 1 1 1 0 1 1 1 1 1 1 1 1 1 1 1 0]

```

## ▼ Evaluating the model and training the Model

Double-click (or enter) to edit

```

model = LogisticRegression()
model.fit(x_train,y_train)
predictions = model.predict(x_test)
print(accuracy_score(y_test, predictions))

```

```
↳ 0.7708333333333334
```

Double-click (or enter) to edit

```

model = DecisionTreeClassifier()

```



```
model.fit(x_train,y_train)
predictions = model.predict(x_test)
print(accuracy_score(y_test, predictions))
```

📄 0.6354166666666666

Double-click (or enter) to edit

```
model = RandomForestClassifier(n_estimators=100)
model.fit(x_train,y_train)
predictions = model.predict(x_test)
print(accuracy_score(y_test, predictions))
```

📄 0.7708333333333334